

# Convolutional Neural Networks & Computer Vision with



## TensorFlow

# Where can you get help?

- Follow along with the code



```
So far we've covered the basics of TensorFlow and built a handful of models to work across different problems. Now we're going to get specific and see how a special kind of neural network, convolutional neural networks \(CNNs\) can be used for computer vision (detecting patterns in visual data).  
Note: In deep learning, many different kinds of model architectures can be used for different problems. For example, you could use a convolutional neural network for making predictions on image data and/or text data. However, in practice some architectures typically work better than others.  
For example, you might want to:

- Classify whether a picture of food contains pizza or steak (we're actually going to do this)
- Detect whether or not an object appears in an image (e.g. did a specific car pass through a security camera?)

In this notebook, we're going to follow the TensorFlow modelling workflow we've been following so far whilst learning about how to build and use CNNs.
```

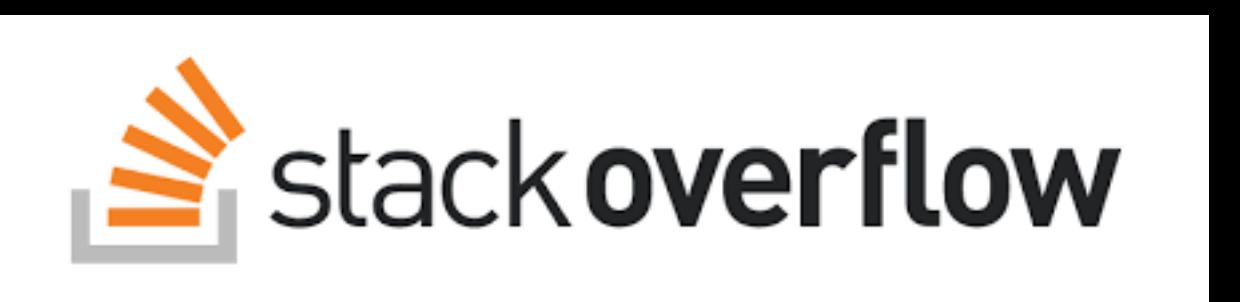
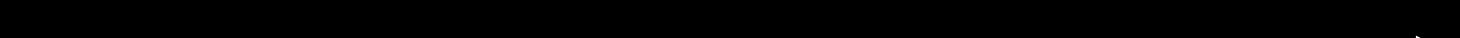
"If in doubt, run the code"

- Try it for yourself



```
# Create a CNN model (same as before)  
model_1 = tf.keras.models.Sequential()  
tf.keras.layers.Conv2D(32, 3, activation='relu'),  
tf.keras.layers.MaxPool2D(pool_size=2, # pool_size can also be (2, 2)  
padding='valid'), # padding can also be 'same'  
tf.keras.layers.Conv2D(32, 3, activation='relu'),  
tf.keras.layers.Conv2D(32, 3, activation='relu'), # activation='relu' == tf.keras.layers.Activations(tf.nn.relu)  
tf.keras.layers.Flatten(),  
tf.keras.layers.Dense(1, activation='sigmoid') # binary activation output  
]  
  
# Compile the model  
model_1.compile(loss='binary_crossentropy',  
optimizer=tf.keras.optimizers.Adam(),  
metrics=['accuracy'])  
  
# Fit the model  
history_1 = model_1.fit(train_data,  
epochs=5,
```

- Press SHIFT + CMD + SPACE to read the docstring

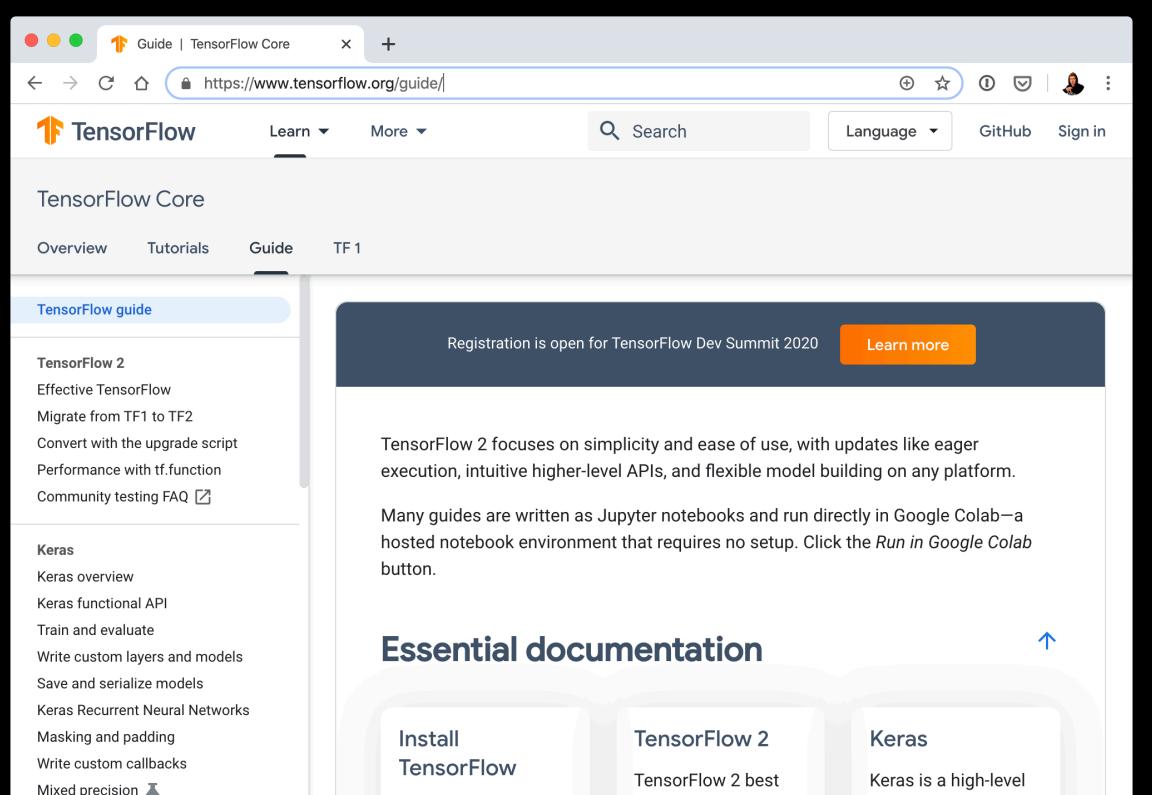


- Search for it

- Try again

- Ask (don't forget the Discord chat!)

(yes, including the "dumb" questions)



“What is a computer vision  
problem?”

# Example computer vision problems



“Is this a photo of sushi, steak or pizza?”

**Binary classification**

(one thing or another)

**Multiclass classification**

(more than one thing or  
another)

**Object detection**

(where's the thing we're  
looking for?)

# What we're going to cover

(broadly)

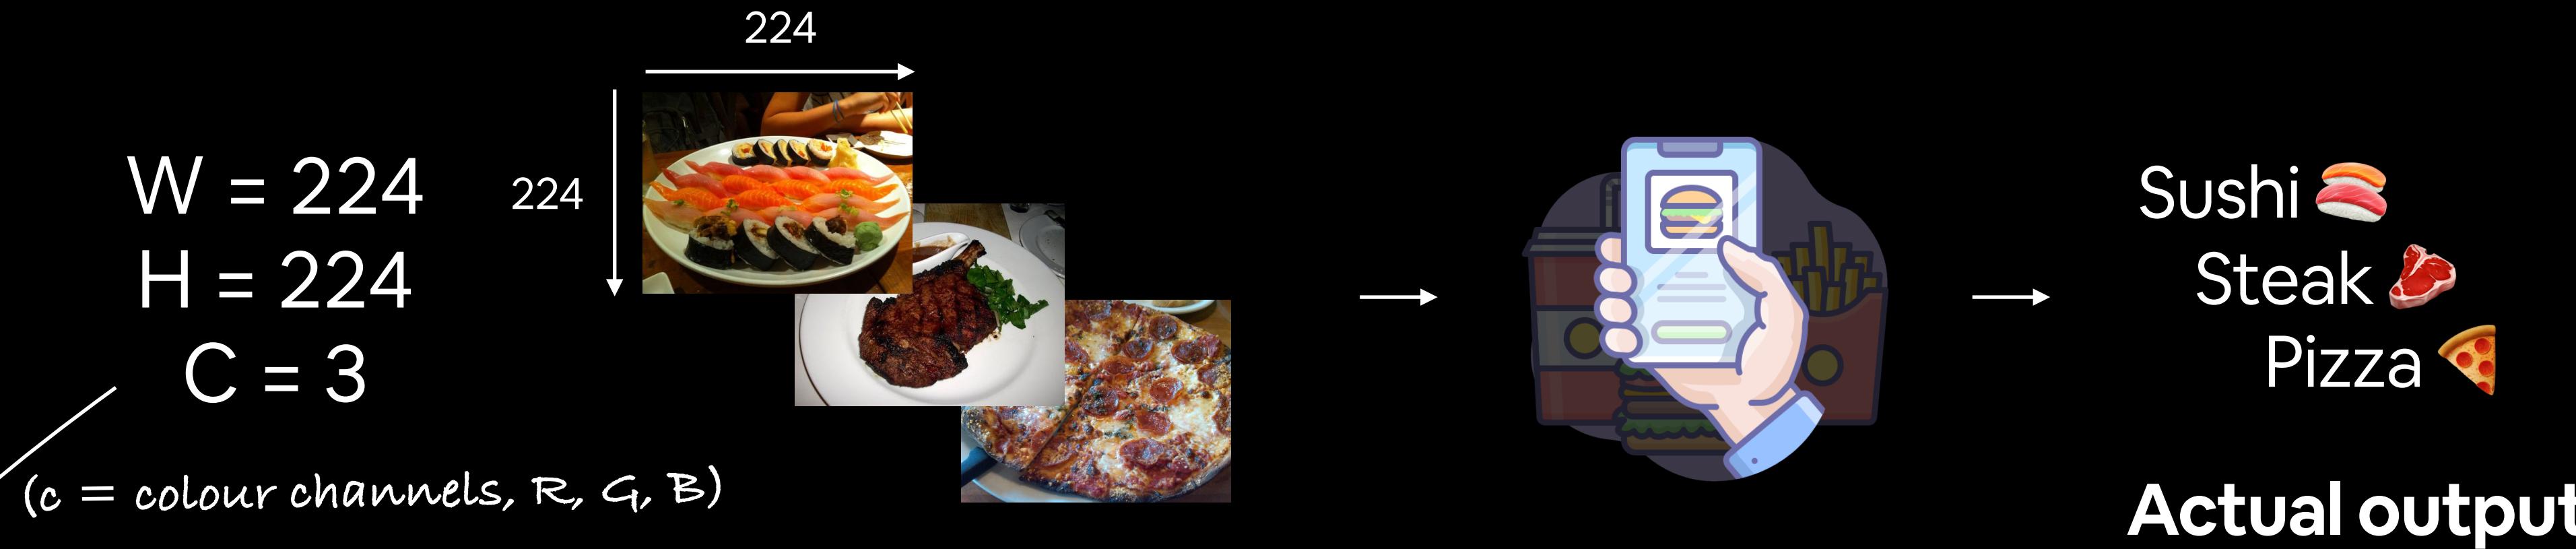
- Getting a dataset to work with (pizza\_steak 🍕🥩)
- Architecture of a **convolutional neural network (CNN)** with TensorFlow
- An end-to-end binary image classification problem
- Steps in modelling with **CNNs**
  - Creating a CNN, compiling a model, fitting a model, evaluating a model
- An end-to-end multi-class image classification problem
- Making predictions on our own custom images

(we'll be cooking up lots of code!)

How:



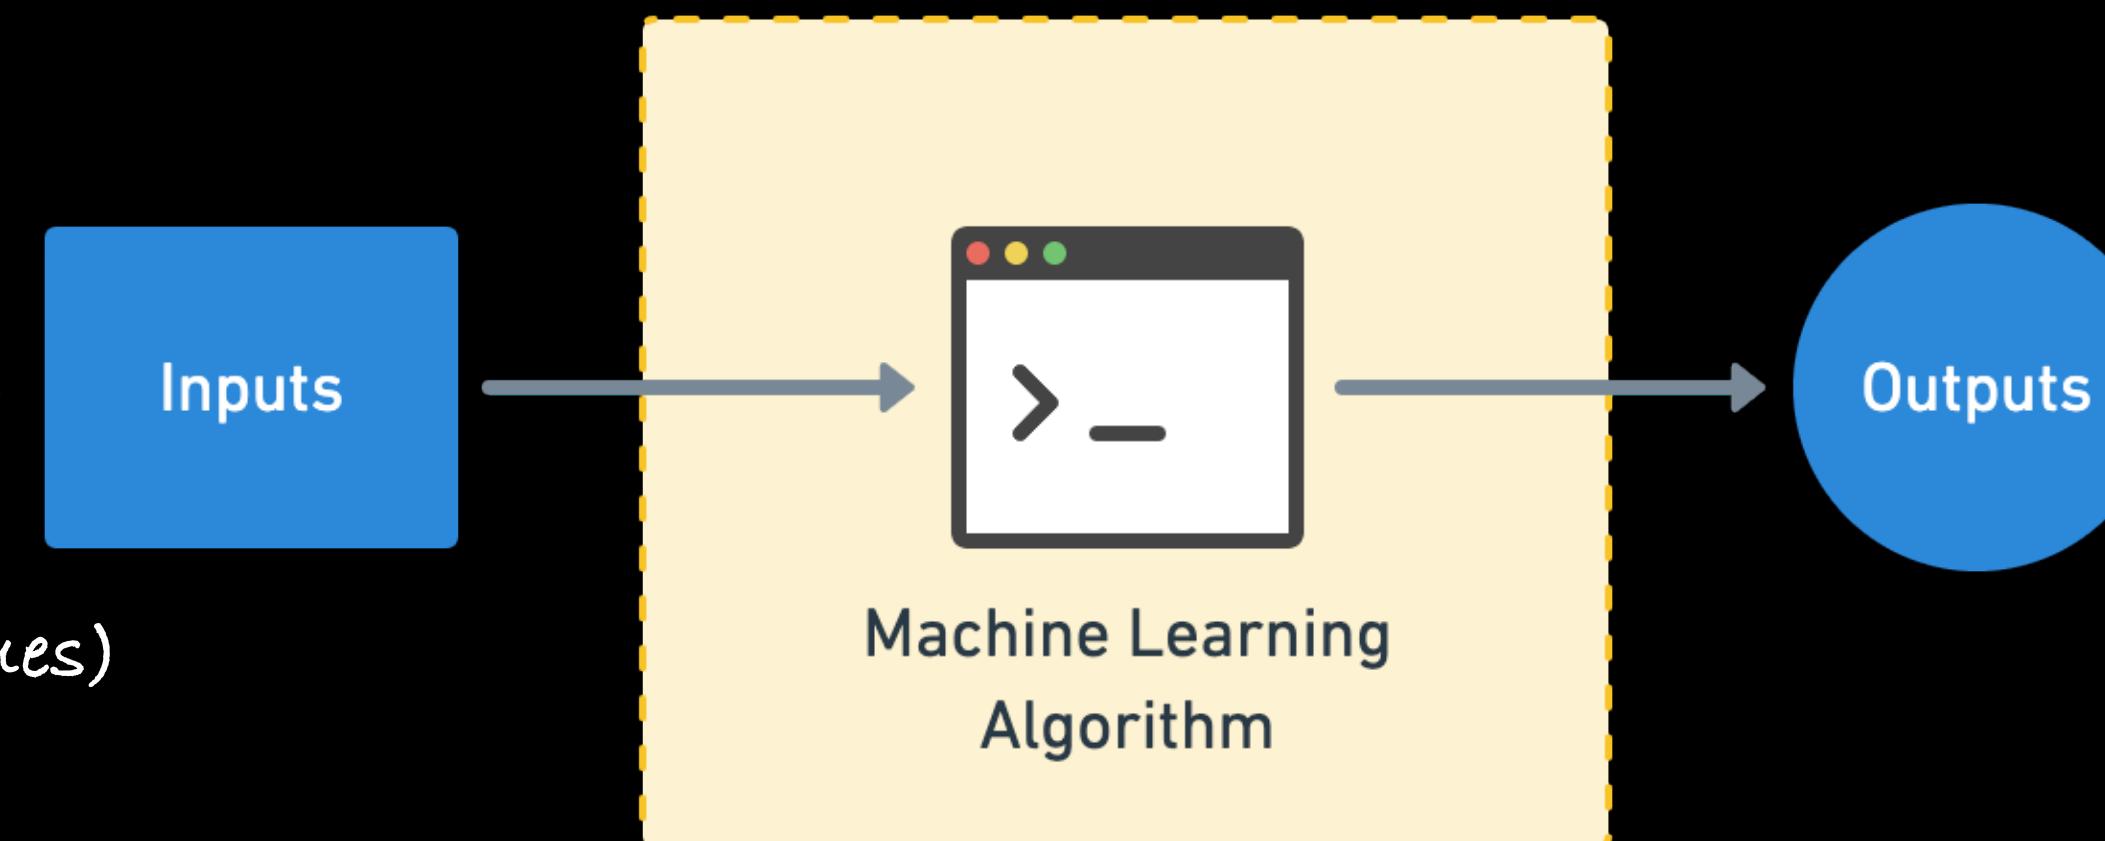
# Computer vision inputs and outputs



This is often a convolutional neural network (CNN)!

$[[0.31, 0.62, 0.44...],$   
 $[0.92, 0.03, 0.27...], \rightarrow$   
 $[0.25, 0.78, 0.07...],$   
 $\dots,$  (normalized pixel values)

Numerical encoding



(often already exists, if not,  
you can build one)

Outputs → Predicted output  
[[0.97, 0.00, 0.03], ✓  
[0.81, 0.14, 0.05], ✗  
[0.03, 0.07, 0.90], ✓  
 $\dots,$

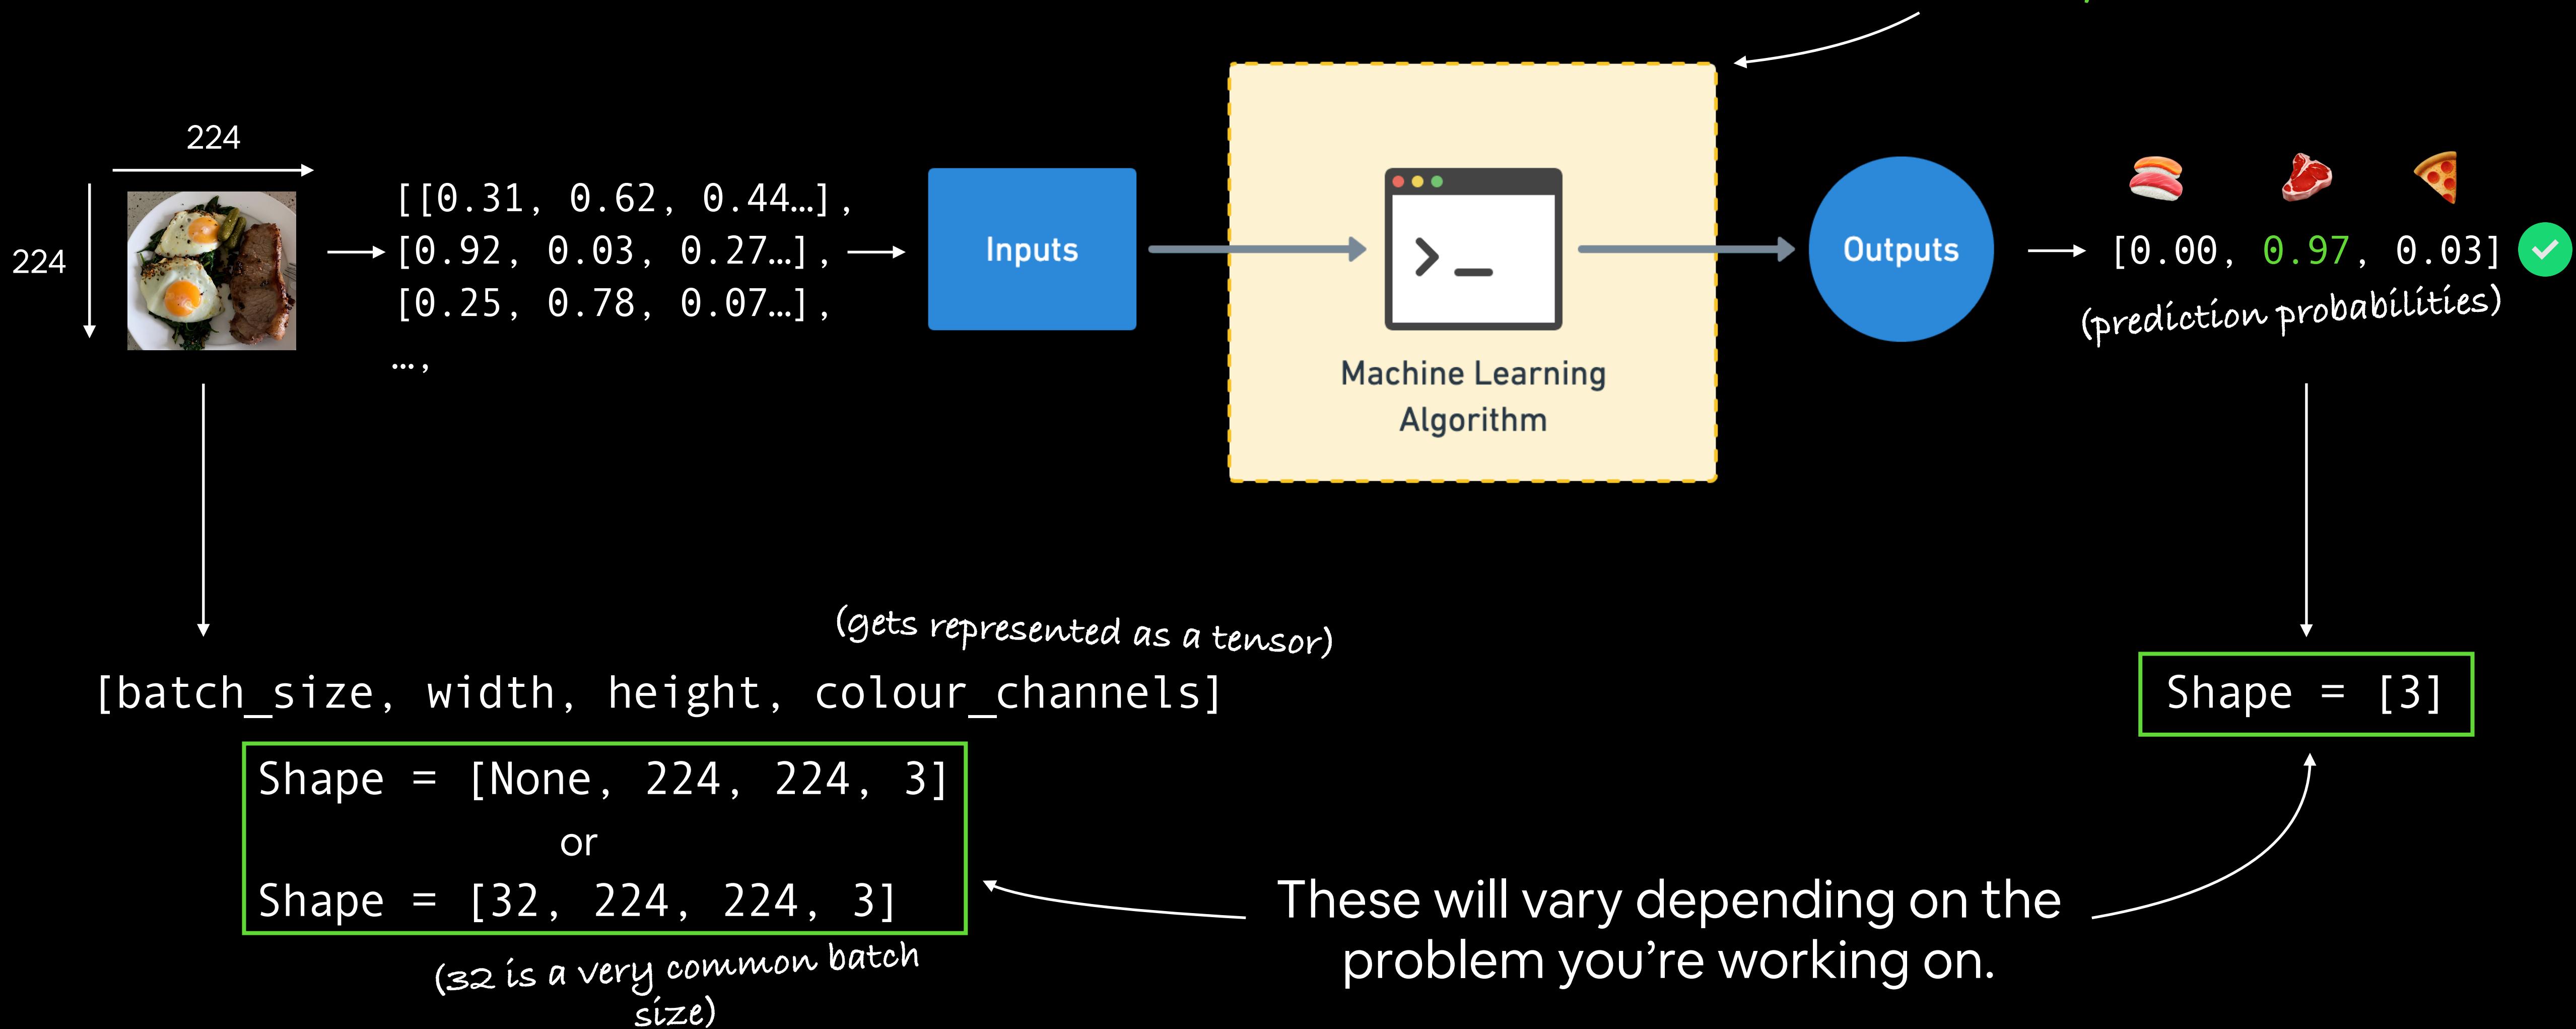
Predicted output

(comes from looking at lots  
of these)

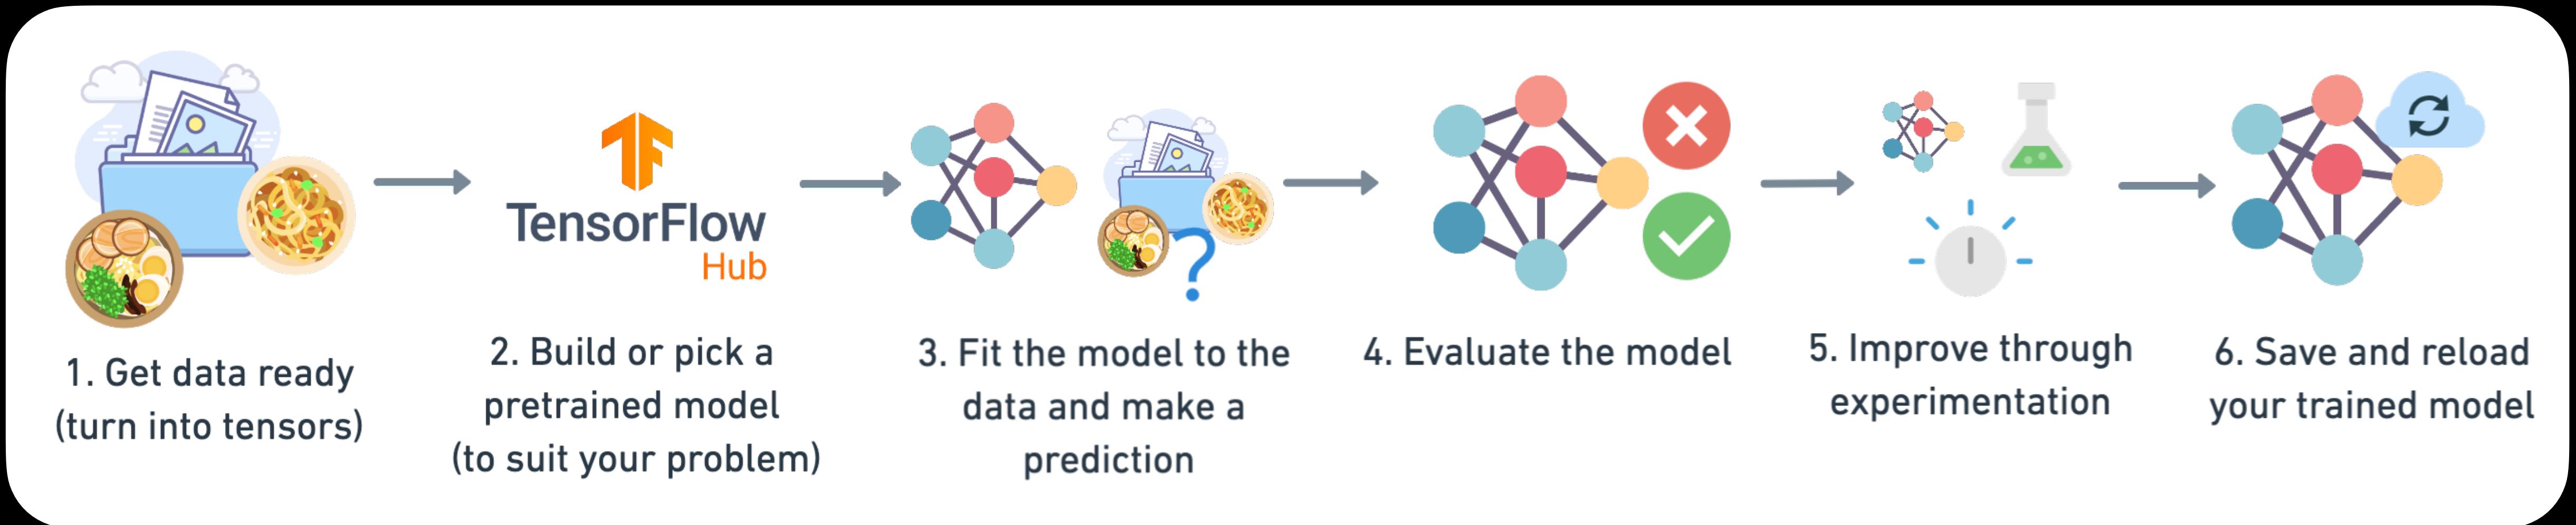
# Input and output shapes

(for an image classification example)

We're going to be building CNNs  
to do this part!



# Steps in modelling with TensorFlow



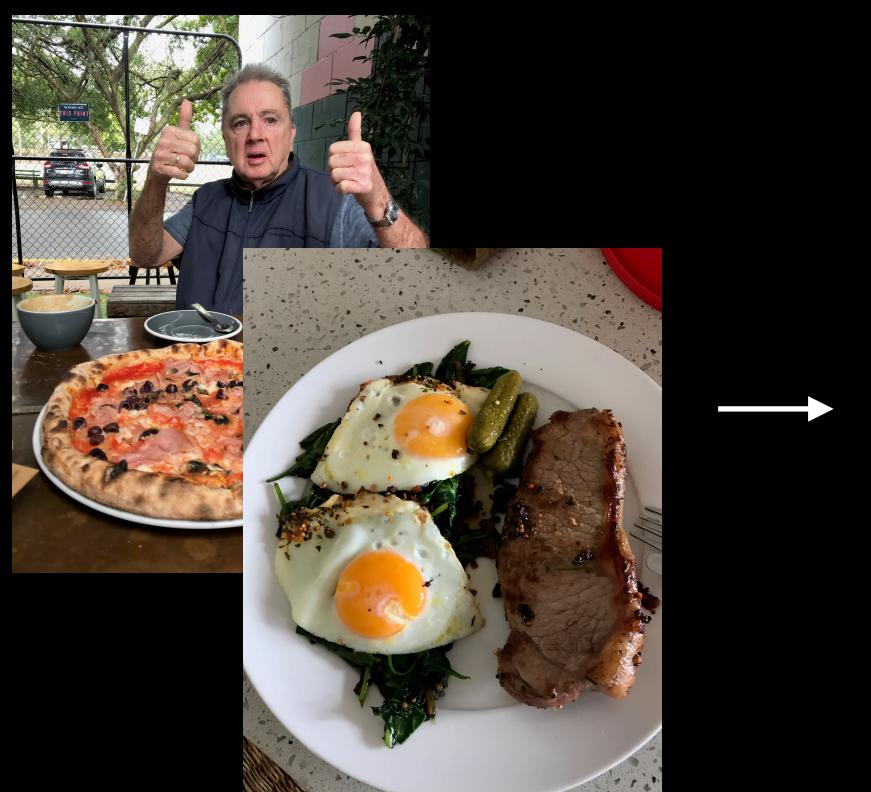
- ↓
1. Turn all data into numbers (neural networks can't handle images)
  2. Make sure all of your tensors are the right shape
  3. Scale features (normalize or standardize, neural networks tend to prefer normalization)

“What is a convolutional neural network (CNN)?”

(typical)\*

# Architecture of a CNN

Hyperparameter/Layer type	What does it do?	Typical values
Input image(s)	Target images you'd like to discover patterns in	Whatever you can take a photo (or video) of
Input layer	Takes in target images and preprocesses them for further layers	<code>input_shape = [batch_size, image_height, image_width, color_channels]</code>
Convolution layer	Extracts/learns the most important features from target images	Multiple, can create with <code>tf.keras.layers.Conv2D</code> (X can be multiple values)
Hidden activation	Adds non-linearity to learned features (non-straight lines)	Usually ReLU ( <code>tf.keras.activations.relu</code> )
Pooling layer	Reduces the dimensionality of learned image features	Average ( <code>tf.keras.layers.AvgPool2D</code> ) or Max ( <code>tf.keras.layers.MaxPool2D</code> )
Fully connected layer	Further refines learned features from convolution layers	<code>tf.keras.layers.Dense</code>
Output layer	Takes learned features and outputs them in shape of target labels	<code>output_shape = [number_of_classes]</code> (e.g. 3 for pizza, steak or sushi)
Output activation	Adds non-linearities to output layer	<code>tf.keras.activations.sigmoid</code> (binary classification) or <code>tf.keras.activations.softmax</code>



```
# 1. Create a CNN model (same as Tiny VGG - https://poloclub.github.io/cnn-explainer/)
cnn_model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(filters=10,
                          kernel_size=3, # can also be (3, 3)
                          activation="relu",
                          input_shape=(224, 224, 3)), # specify input shape (height, width, colour channels)
    tf.keras.layers.Conv2D(10, 3, activation="relu"),
    tf.keras.layers.MaxPool2D(pool_size=2, # pool_size can also be (2, 2)
                           padding="valid"), # padding can also be 'same'
    tf.keras.layers.Conv2D(10, 3, activation="relu"),
    tf.keras.layers.Conv2D(10, 3, activation="relu"), # activation='relu' == tf.keras.layers.Activations(tf.nn.relu)
    tf.keras.layers.MaxPool2D(2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(1, activation="sigmoid") # binary activation output
])

# 2. Compile the model
cnn_model.compile(loss="binary_crossentropy",
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=[ "accuracy" ])

# 3. Fit the model
history = cnn_model.fit(train_data, epochs=5)
```

(what we're working towards building)

→ Steak Pizza

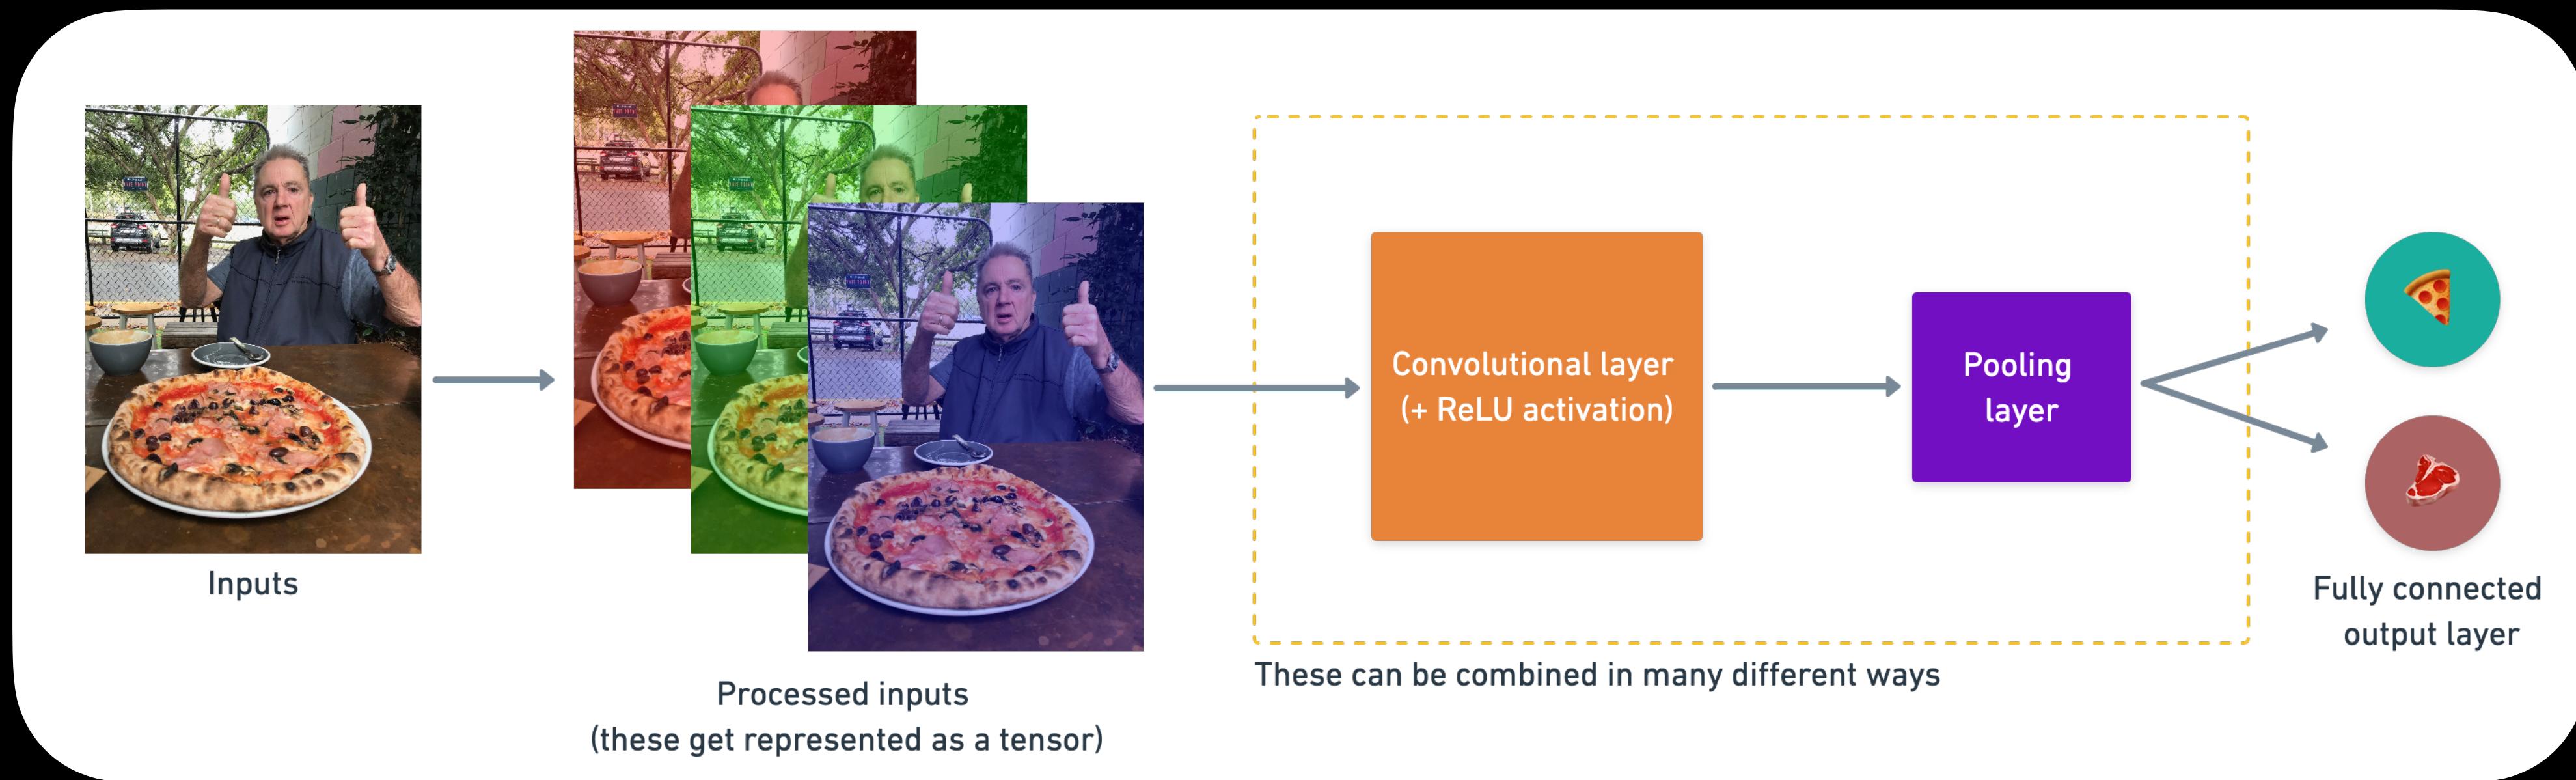
\*Note: there are almost an unlimited amount of ways you could stack together a convolutional neural network, this slide demonstrates only one.

Let's code!

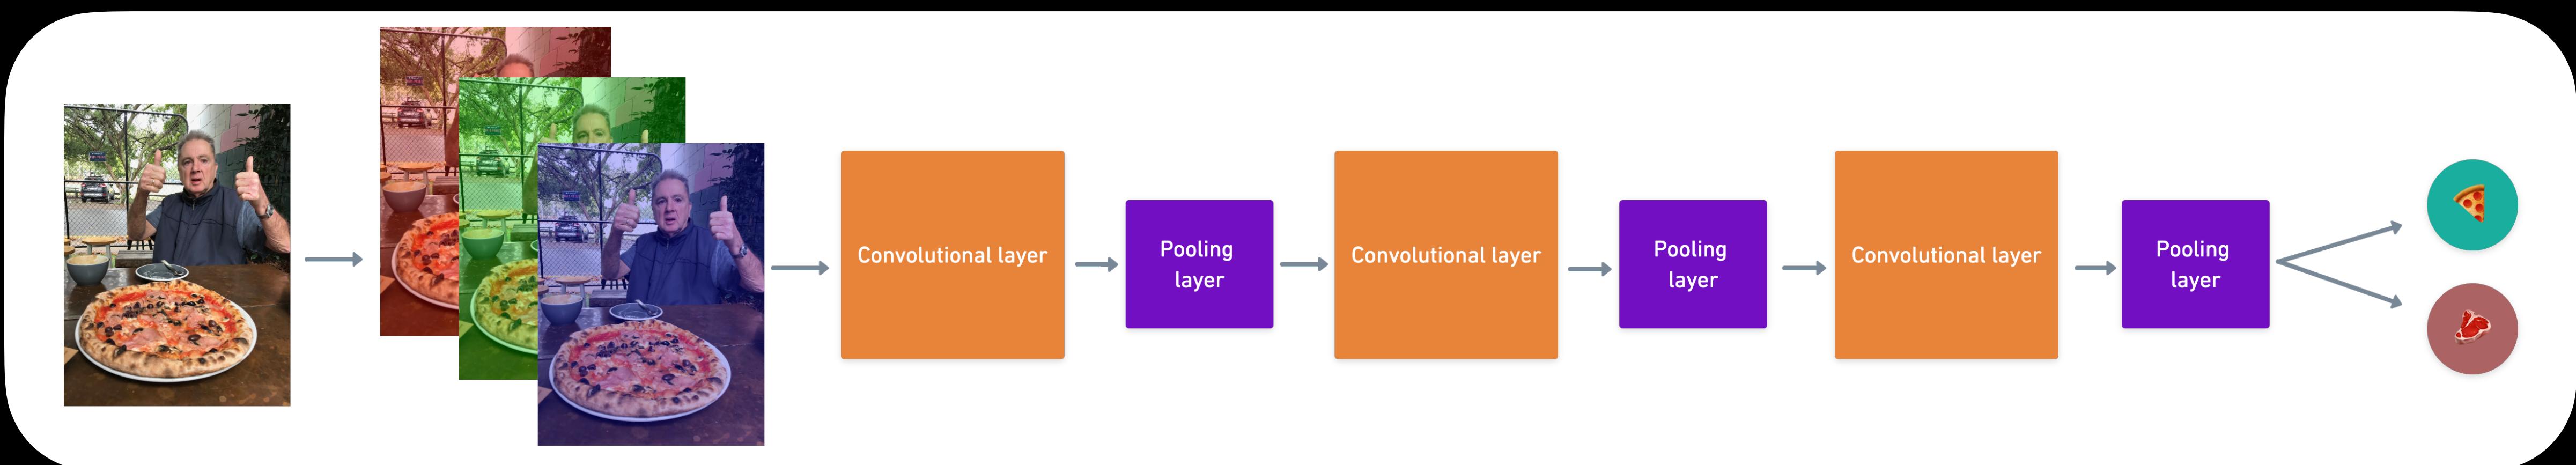
# Architecture of a CNN

(coloured block edition)

## Simple CNN



## Deeper CNN



# Breakdown of Conv2D layer

**Example code:** `tf.keras.layers.Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding="same")`

**Example 2 (same as above):** `tf.keras.layers.Conv2D(filters=10, kernel_size=3, strides=1, padding="same")`

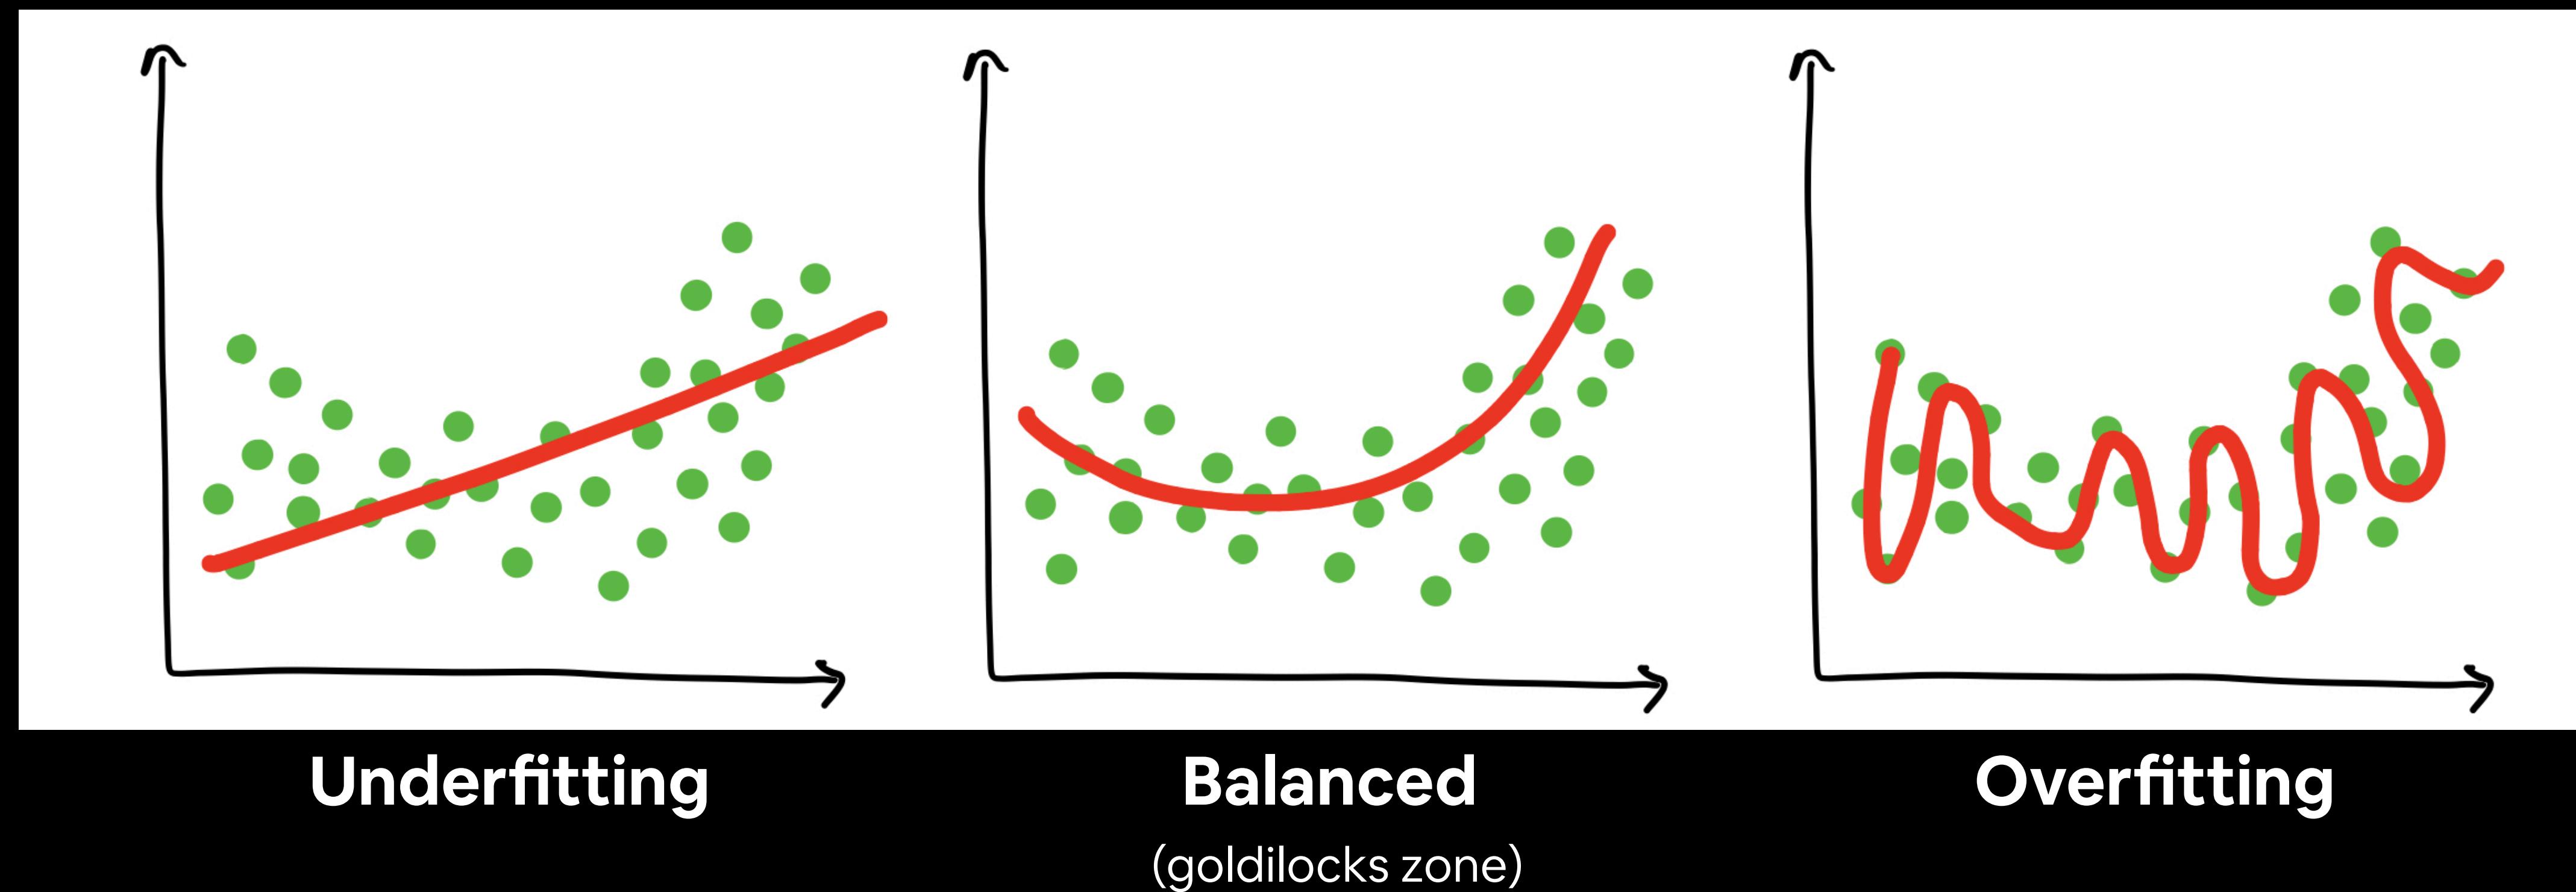
Hyperparameter name	What does it do?	Typical values
Filters	Decides how many filters should pass over an input tensor (e.g. sliding windows over an image).	10, 32, 64, 128 (higher values lead to more complex models)
Kernel size (also called filter size)	Determines the shape of the filters (sliding windows) over the output.	3, 5, 7 (lower values learn smaller features, higher values learn larger features)
Padding	Pads the target tensor with zeroes (if "same") to preserve input shape. Or leaves in the target tensor as is (if "valid"), lowering output shape.	"same" or "valid"
Strides	The number of steps a filter takes across an image at a time (e.g. if <code>strides=1</code> , a filter moves across an image 1 pixel at a time).	1 (default), 2

 **Resource:** For an interactive demonstration of the above hyperparameters, see the [CNN explainer website](#).

# What is overfitting?

**Overfitting** — when a model over learns patterns in a particular dataset and isn't able to generalise to unseen data.

For example, a student who studies the course materials too hard and then isn't able to perform well on the final exam. Or tries to put their knowledge into practice at the workplace and finds what they learned has nothing to do with the real world.



# Improving a model

(from a model's perspective)

```
# 1. Create the model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

# 2. Compile the model
model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(lr=0.001),
              metrics=[ "accuracy" ])

# 3. Fit the model
model.fit(x_train_subset, y_train_subset, epochs=5)
```

Smaller model

```
# 1. Create the model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

# 2. Compile the model
model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(lr=0.0001),
              metrics=[ "accuracy" ])

# 3. Fit the model
model.fit(x_train_full, y_train_full, epochs=100)
```

Larger model

## Common ways to improve a deep model:

- Adding layers
- Increase the number of hidden units
- Change the activation functions
- Change the optimization function
- Change the learning rate      (because you can alter each of these, they're hyperparameters)
- Fitting on more data
- Fitting for longer

# Improving a model

(from a data perspective)

## Method to improve a model (reduce overfitting)

### What does it do?

More data

Gives a model more of a chance to learn patterns between samples (e.g. if a model is performing poorly on images of pizza, show it more images of pizza).

Data augmentation

Increase the diversity of your training dataset without collecting more data (e.g. take your photos of pizza and randomly rotate them 30°). Increased diversity forces a model to learn more generalisation patterns.

Better data

Not all data samples are created equally. Removing poor samples from or adding better samples to your dataset can improve your model's performance.

Use transfer learning

Take a model's pre-learned patterns from one problem and tweak them to suit your own problem. For example, take a model trained on pictures of cars to recognise pictures of trucks.

# What is data augmentation?

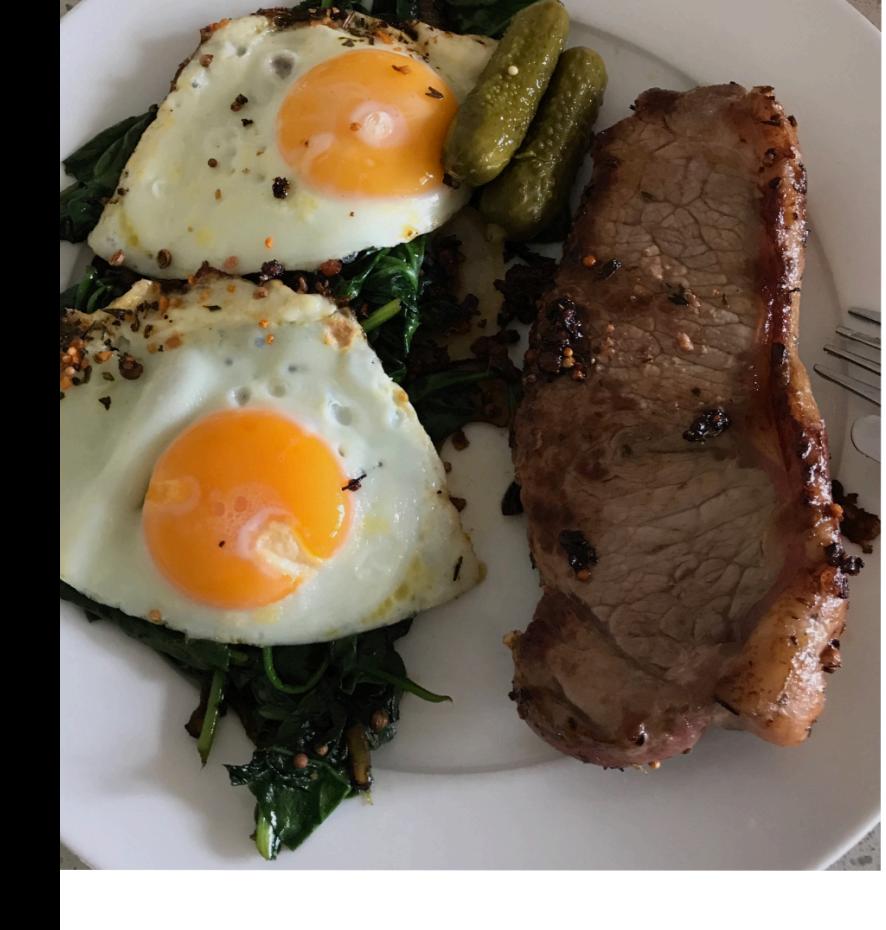
Looking at the same image but from different perspective(s)\*.



Original



Rotate



Shift



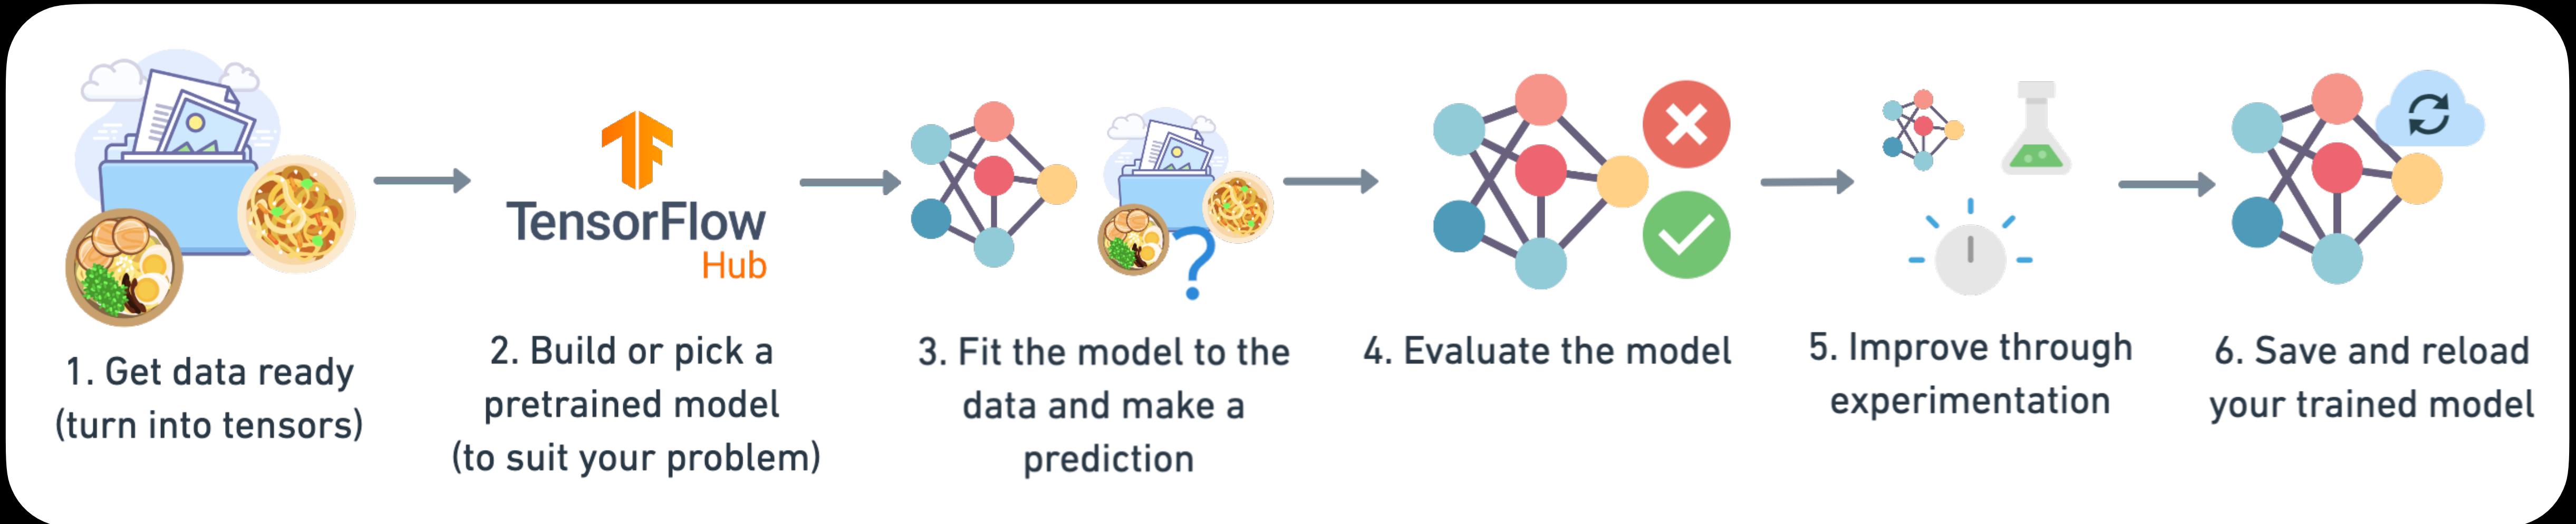
Zoom

\*Note: There are many more different kinds of data augmentation such as, cropping, replacing, shearing. This slide only demonstrates a few.

# Popular & useful computer vision architectures

Architecture	Release Date	Paper	Use in TensorFlow	When to use
ResNet (residual networks)	2015	<a href="https://arxiv.org/abs/1512.03385">https://arxiv.org/abs/1512.03385</a>	Find pre-trained versions on <a href="#">TensorFlow Hub</a> or <a href="#">tf.keras.applications</a>	A good backbone for many computer vision problems
EfficientNet(s)	2019	<a href="https://arxiv.org/abs/1905.11946">https://arxiv.org/abs/1905.11946</a>	Find pre-trained versions on <a href="#">TensorFlow Hub</a> or <a href="#">tf.keras.applications</a>	Typically now better than ResNets for computer vision
MobileNet(s)	2017	<a href="https://arxiv.org/abs/1704.04861">https://arxiv.org/abs/1704.04861</a>	Find pre-trained versions on <a href="#">TensorFlow Hub</a> or <a href="#">tf.keras.applications</a>	Lightweight architecture suitable for devices with less computing power

# Steps in modelling with TensorFlow



```
# 1. Create a model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.Input(shape=(224, 224, 3)),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(3, activation="softmax")
])

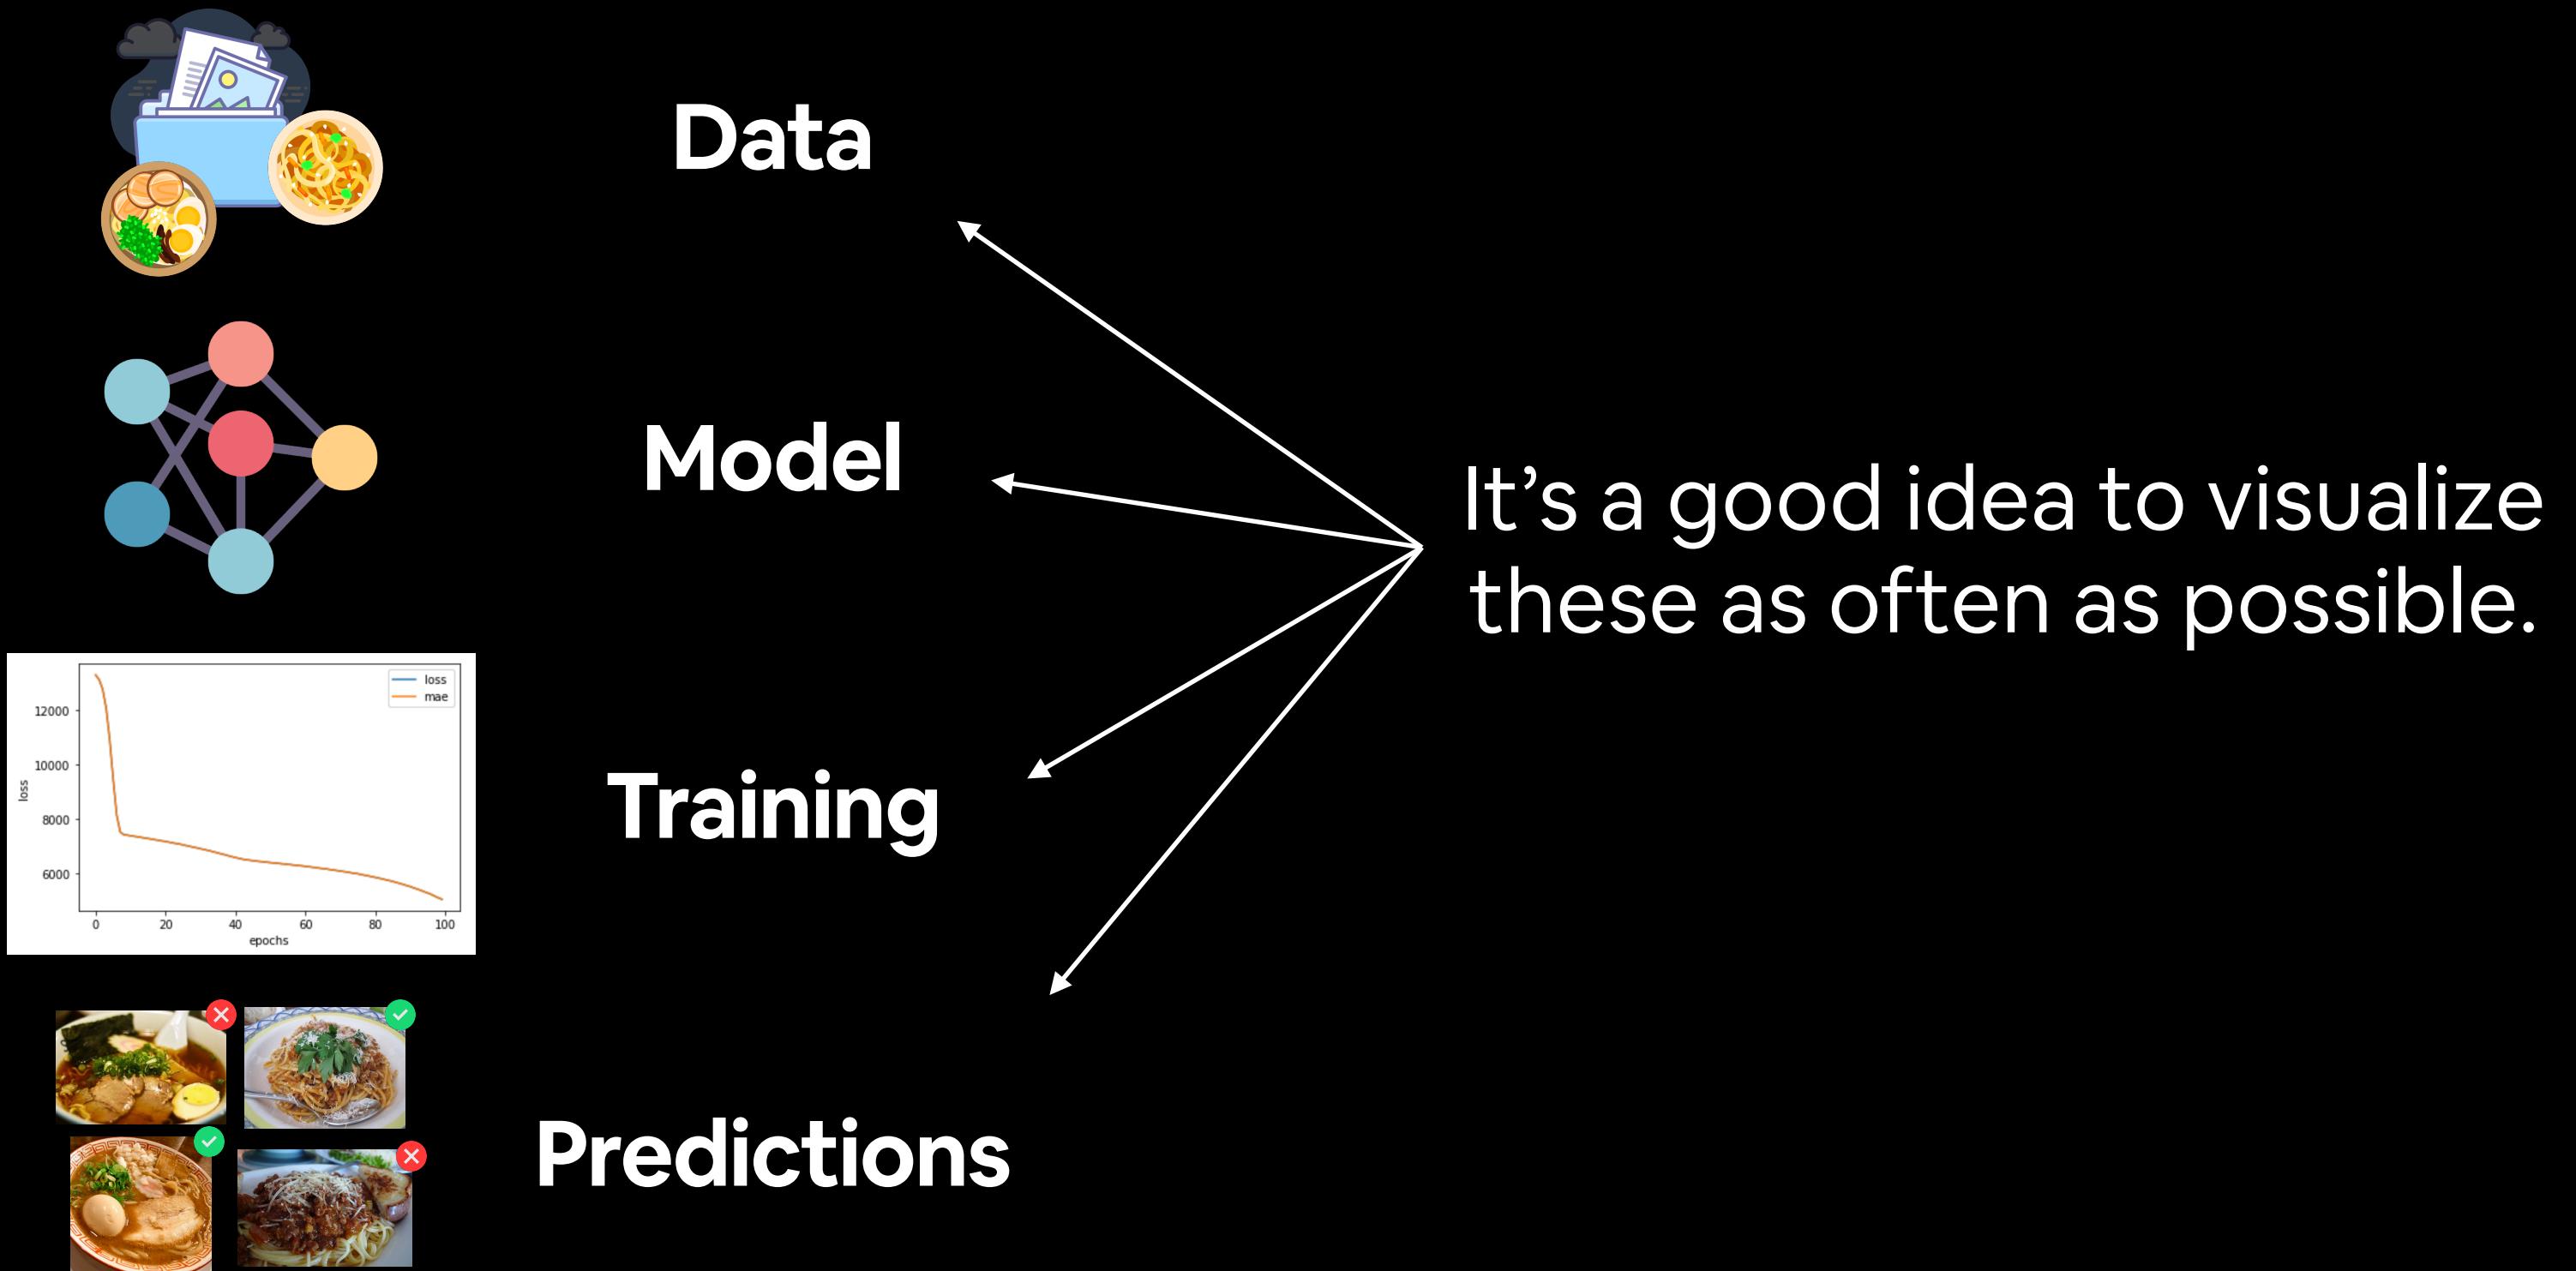
# 2. Compile the model
model.compile(loss=tf.keras.losses.CategoricalCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(),
              metrics=[ "accuracy" ])

# 3. Fit the model
model.fit(X_train, y_train, epochs=5)

# 4. Evaluate the model
model.evaluate(X_test, y_test)
```

# The machine learning explorer's motto

“Visualize, visualize, visualize”



# The machine learning practitioner's motto

“Experiment, experiment, experiment”



*(try lots of things and see what  
tastes good)*